

Welcome to **Day 1** of our **SRE Database Training Module (v3.0)**! Today, we'll explore **Relational Database Fundamentals & Basic SQL Queries** and lay a solid foundation for your journey toward operational excellence. We'll focus on:

- How databases store data in tables, columns, and rows
- Key relational concepts such as primary keys and foreign keys
- Using **SELECT**, **FROM**, and **WHERE** to retrieve data efficiently
- How these fundamentals appear in Oracle, PostgreSQL, SQL Server, MongoDB (NoSQL), and Kafka (streaming)
- Common pitfalls, career protection strategies, and SRE-oriented best practices

Databases are the backbone of reliable applications—poor query practices can cause performance degradation, data inconsistencies, or even system outages. Real incidents have seen entire production systems brought down by a single careless query. **Why does this matter?** Because as a support engineer or SRE, having a firm grasp on data retrieval and relational design can drastically cut downtime, speed up issue resolution, and protect your organization from costly mistakes.

Below is a high-level **concept map** of what we'll cover:

Throughout this module, we'll build each concept **brick by brick**, ensuring each step is logically connected to the next. Let's dive in!

& Learning Objectives by Tier

Beginner Objectives (4)

- 1. **Identify** tables, columns, and rows in a relational database, explaining them in simple real-world analogies.
- 2. **Explain** the roles of primary keys and foreign keys in maintaining data integrity.
- 3. Construct basic SQL queries using SELECT, FROM, and WHERE to retrieve specific data.
- 4. **Connect** to sample databases and perform fundamental data lookups without errors.



- 1. **Analyze** table relationships and how primary/foreign keys impact support guery design.
- 2. **Optimize** basic queries by choosing relevant columns and applying WHERE clauses efficiently.
- 3. Compare different database systems (Oracle, PostgreSQL, SQL Server, MongoDB, Kafka) for data retrieval needs.
- 4. **Apply** early SRE principles—like observability and minor performance checks—to daily queries.

SRE-Level Objectives (4)

- 1. **Evaluate** how table structure and keys affect long-term reliability and performance under load.
- 2. **Investigate** slow or failing queries using preliminary execution plan analysis.
- 3. **Translate** gueries and concepts across relational, NoSQL, and streaming platforms as needed.
- 4. Implement foundational reliability practices (e.g., indexing, basic monitoring) for critical queries.



Knowledge Bridge

Before we start:

- Prerequisite Knowledge: Familiarity with basic computing concepts (files, processes, basic OS usage).
- Connections to Prior Knowledge: If you've ever used a spreadsheet, you already understand columns and rows. We're just applying more formal rules and constraints.
- Foundation for Future Topics: Mastery of Day 1 content is critical for more advanced operations like INSERT, UPDATE, DELETE, and performance tuning.
- Cross-Paradigm Link: Even if you deal primarily with NoSQL or streaming data, relational principles often guide how you structure your data or logs.

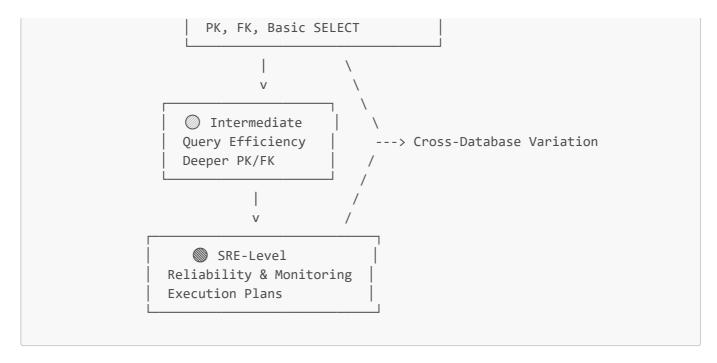
Learning Journey Timeline:

```
[ Day 1: Relational Fundamentals & Basic SQL ]
[ Day 2: CRUD + Transactions ] -> [ Later: Joins, Indexing, Tuning ] -> [ SRE-
level HA & Observability ]
```

■ Visual Concept Map

Below is a more detailed visual concept map showing how today's topics fit together. The color-coded blocks indicate complexity tiers:

```
Fundamentals
Tables, Columns, Rows
```



- Relational Concepts anchor all tiers.
- Each tier focuses on a new depth of detail and SRE integration.
- Cross-database comparisons apply at every level.

Core Concepts

1. Relational Database Fundamentals

- **Beginner Analogy**: Think of a **digital filing cabinet** where each drawer is a **table**, each folder a **row**, and each labeled section a **column**.
- **W** Visual Representation:



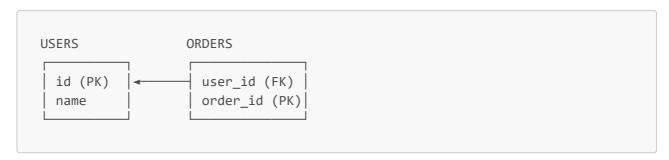
- <u>A</u> **Technical Explanation**: Data is logically grouped into tables. Each table has a set of columns (fields) that define the data attributes, and rows (records) store actual data.
- Support/SRE Application: Daily tasks often include looking up user records or referencing error logs. Understanding table structure is vital.
- System Impact: Well-structured tables enhance data consistency; poor design can lead to duplicates or anomalies.

 Oracle / PostgreSQL / SQL Server: All rely on the core relational model, though syntax for table creation differs slightly.

- MongoDB: Doesn't use tables/rows; uses collections (like tables) and documents (like rows).
- Kafka: Data is stored in topics. The idea of "rows" might map to a "message" in a topic.
- **Common Misconceptions**: Thinking all databases are relational. Some handle data differently (MongoDB, Kafka).
- Quick Reference: "A table is a structured container for records with uniform column definitions."

2. Key Concepts: Primary & Foreign Keys

- **Beginner Analogy**: A **primary key** is like a **unique driver's license number** for each person. A **foreign key** is like referencing that license number in a traffic database to track violations.
- **W** Visual Representation:



- <u>\$\Delta\$</u> Technical Explanation:
 - o Primary Key (PK): Uniquely identifies each row. No duplicates allowed.
 - Foreign Key (FK): References a PK in another table, ensuring relational integrity.
- **Support/SRE Application**: When searching for user transactions, a foreign key ensures you correctly link the user to the relevant orders table.
- System Impact: Proper PK/FK usage ensures consistent data retrieval, easier debugging, and eliminates orphaned records.
- - Oracle / PostgreSQL / SQL Server: Typically enforce PK/FK constraints by default if defined.
 - MongoDB: No strict FKs; references might be manual or embedded documents.
 - **Kafka**: No concept of PK/FK in message streams. Keys can be used for partitioning, but not relational integrity.
- **Common Misconceptions**: Some assume foreign keys are automatically created or enforced. Must be explicitly defined in relational systems.
- Quick Reference: "PK = Unique row ID. FK = Link to that ID in another table."

3. Basic SQL: SELECT, FROM, WHERE

- **Beginner Analogy**: "Asking a librarian" for specific books (SELECT) from a certain shelf (FROM) that match a certain topic (WHERE).
- Wisual Representation:

```
SELECT columns
FROM table
WHERE condition;
```

```
Database Engine
(parses & executes SQL)
returns matching rows
```

- <u>\$\Delta\$</u> Technical Explanation:
 - SELECT: Chooses which columns to display
 - FROM: Specifies which table(s) or data source(s) to query
 - WHERE: Filters rows based on specified conditions
- **Support/SRE Application**: Useful for triaging tickets ("Which users are affected by this bug?") or verifying data states.
- System Impact: Well-written queries prevent excessive CPU/I/O. Missing WHERE clauses can return huge results, straining the system.
- - Oracle / PostgreSQL / SQL Server: Similar SQL syntax with minor differences (e.g., quoting identifiers).
 - MongoDB: db.collection.find({ condition }) instead of SELECT/FROM/WHERE.
 - Kafka: Queries might be done with KSQL or by consuming messages from topics.
- **Common Misconceptions**: Overusing SELECT * is harmless. In production, it can degrade performance drastically.
- Quick Reference: "SQL 'SELECT-FROM-WHERE' is the fundamental pattern for retrieving data in relational databases."



Day 1 Concept & Command Breakdown

Command/Concept: Relational Database Structure

Overview:

Defines how data is organized into tables, columns, and rows within a schema. Essential for understanding any subsequent SQL operation.

Real-World Analogy:

A binder with dividers (tables) and pages (rows), each page containing the same sections (columns).

Visual Representation:



Syntax & Variations:

Syntax Form	Example	Description	Support/SRE Usage Context
Create Table	CREATE TABLE users (id INT PRIMARY KEY, name VARCHAR(50));	Defines a new table structure.	Building or modifying schema in dev/test
Basic Table Query	SELECT * FROM users;	Retrieves all rows/columns.	Quick inspection of stored data
Advanced Usage	SELECT COUNT(*) FROM users;	Summarizes table data.	Checking user volume for capacity planning

Cross-Database Implementation:

Database System	Syntax/Implementation	Special Features	Limitations
Oracle	Uses CREATE TABLE with Oracle-specific data types (e.g., NUMBER).	Partitioned tables for large- scale enterprise usage.	Some older versions have strict naming.
PostgreSQL	Very ANSI-compliant, can store JSON columns for semi-structured data.	Rich indexing capabilities (GIN, BRIN).	May require tuning for very large DBs.
SQL Server	T-SQL with CREATE TABLE Offers identity columns for PK.	Tight integration with Microsoft ecosystem (SSMS).	Windows-based licensing considerations.
MongoDB	Collections created implicitly. Documents can vary in structure.	Flexible schema, no formal table definition.	No direct enforcement of schema.
Kafka	No tables, but topics. Concept of data "structure" is in message format.	Streaming logs can approximate table-like queries with KSQL.	Not strictly relational.

Tiered Examples:


```
-- Example: Show all columns/rows from 'users'

SELECT *

FROM users;

/* Expected Output:

| id | name |

|----|-----|
```

```
| 1 | Bob | | 2 | Alice | ... */
-- Step-by-step:
-- 1) The query references the table "users"
-- 2) SELECT * retrieves all columns
-- 3) No WHERE means get every row
```



```
-- Example: Summarize how many rows in 'orders'

SELECT COUNT(*) AS total_orders

FROM orders;

/* Expected Output:

| total_orders |

|------|

| 42 |

*/

-- Support Relevance:
-- 1) Quick check to confirm how many orders exist
-- 2) Used for capacity or sanity checks
```

• SRE-Level Example:

Instructional Notes:

- @ Beginner Tip: Always know how to list tables (e.g., SHOW TABLES; in MySQL) to avoid referencing nonexistent objects.
- Beginner Tip: Understand your environment's naming conventions. Some DBs require quotes or brackets.

• SRE Insight: Monitoring table growth is critical to preventing surprises in storage usage.

- SRE Insight: Large "hot" tables might need partitioning or archiving strategies.
- **Common Pitfall:** Using ambiguous table/column names leads to confusion (e.g., naming a column "date").
- **Common Pitfall:** Creating tables without a primary key often leads to duplicates or messy data.
- **Security Note:** Restrict CREATE/ALTER table privileges to authorized roles.
- **Performance Impact:** Proper table design can drastically reduce query times.
- Career Risk: Accidentally dropping or modifying the schema in production can cause immediate incidents.
- Recovery Strategy: Ensure up-to-date backups or point-in-time recovery are configured before schema changes.

Command/Concept: Primary Keys & Foreign Keys

Overview:

Primary keys guarantee uniqueness, while foreign keys establish relationships between tables. They're critical for relational integrity.

Real-World Analogy:

PK = a unique ID card. **FK** = referencing that ID in another system to verify the same person.

Visual Representation:



Syntax & Variations:

Syntax			Support/SRE
Form	Example	Description	Usage
roilli			Context

Syntax Form	Example	Description	Support/SRE Usage Context
Create PK	CREATE TABLE users (id INT PRIMARY KEY, name VARCHAR(50));	Single column PK ensures each id is unique.	Basic table creation.
Create FK	ALTER TABLE orders ADD CONSTRAINT fk_user FOREIGN KEY (user_id) REFERENCES users(id);	Links user in orders to users table.	Data integrity enforcement.
Advanced Usage	<pre>CREATE TABLE memberships (user_id, group_id, PRIMARY KEY (user_id, group_id));</pre>	Composite PK for linking multiple keys.	Many-to- many relationships.

Cross-Database Implementation:

Database System	Syntax/Implementation	Special Features	Limitations
Oracle	CONSTRAINT pk_name PRIMARY KEY (col_name) USING INDEX	Rich indexing & partitioning.	Some older versions might require manual index creation for PK.
PostgreSQL	PRIMARY KEY (col_name) or UNIQUE + NOT NULL	Extremely flexible with constraints.	Must watch out for concurrency if not carefully tuned.
SQL Server	PRIMARY KEY CLUSTERED(col_name)	Clustered vs. nonclustered indexes.	Licensing can be expensive for large scale.
MongoDB	_id field is auto-generated PK per document. No real FK constraints.	Flexible referencing strategies.	No enforced relational constraints.
Kafka	"Key" used for partitioning, not relational integrity.	Ideal for streaming partition logic.	Doesn't enforce cross-message constraints.

Tiered Examples:


```
-- Ensure 'id' is PK in a new 'users' table

CREATE TABLE users (
   id INT PRIMARY KEY,
   name VARCHAR(50)
);

/* Output: Table 'users' created with 'id' as PK */
```

```
-- Step-by-step:
-- 1) 'PRIMARY KEY' ensures unique, non-null 'id'
-- 2) Basic building block for relational data
```

• **(a)** Intermediate Example:

```
-- Add a foreign key to link orders to users

ALTER TABLE orders

ADD CONSTRAINT fk_user

FOREIGN KEY (user_id) REFERENCES users(id);

/* Output: Orders now references Users */

-- Support scenario:

-- Prevent orphan orders by guaranteeing user_id must exist in Users
```

• SRE-Level Example:

```
-- Composite primary key to handle many-to-many user-group relationships
CREATE TABLE memberships (
   user_id INT NOT NULL,
   group_id INT NOT NULL,
   PRIMARY KEY (user_id, group_id)
);
/* Output: PK ensures each user-group pair is unique */
-- Production Relevance:
-- Large-scale identity management relies on composite PK for linking multiple entities
```

Instructional Notes:

- SRE Insight: Without correct PK/FK, data anomalies can lead to inaccurate metrics or analytics.
- SRE Insight: Monitoring constraints helps detect unexpected data behavior quickly.
- **Common Pitfall:** Disabling FK constraints for "speed" can cause massive data consistency issues later.
- **Common Pitfall:** Using "natural keys" (like emails) might cause complexity if emails change.
- **Security Note:** Ensuring correct references prevents data exposure issues (orphaned sensitive data).
- Performance Impact: Proper indexing on PK/FK columns speeds up joins and lookups.
- 🙀 Career Risk: Accidentally dropping or altering a PK can cascade foreign key failures across the DB.
- Recovery Strategy: If constraints break, fix schema references or restore from a known good state.

Command/Concept: SELECT, FROM, WHERE

Overview:

The foundational SQL statements for retrieving data. **SELECT** chooses columns, **FROM** picks the table, and **WHERE** filters rows.

Real-World Analogy:

SELECT = "Which details of a library book do I want?"

FROM = "Which shelf of books am I looking at?"

WHERE = "Which genre or author am I focusing on?"

Visual Representation:

Syntax & Variations:

Syntax Form	Example	Description	Support/SRE Usage Context
Basic SELECT	SELECT * FROM users;	Retrieves all columns, all rows.	Quick data check or debugging.
Filtered WHERE	SELECT name FROM users WHERE region='NA';	Only rows where region='NA'.	Focusing on a subset, e.g. region.
Advanced Usage (JOIN)	SELECT u.name, o.order_id FROM users u JOIN orders o ON u.id=o.user_id;	Combines data across multiple tables.	Investigating user orders in one query.

Cross-Database Implementation:

Database System	Syntax/Implementation	Special Features	Limitations
Oracle	Similar to ANSI SQL, often uses double quotes for identifiers	Rich date/time functions.	Quoting needed if object names are cAsE sensitive.
PostgreSQL	Very ANSI-compliant, uses ILIKE for case-insensitive.	JSON support for semi-structured data.	Might require extra config for advanced indexing.

Database System	Syntax/Implementation	Special Features	Limitations
SQL Server	T-SQL variant, [] for quoted identifiers.	Additional T-SQL features like TOP.	Some features differ from pure ANSI SQL.
MongoDB	<pre>db.collection.find({region:'NA'}) instead of SELECT/FROM.</pre>	Flexible queries with embedded docs.	Not truly SQL-based.
Kafka	Use KSQL (e.g., SELECT * FROM topic_stream WHERE)	Real-time stream filtering.	Queries ephemeral, data streaming.

Tiered Examples:


```
-- Retrieve all columns from 'customers'

SELECT *

FROM customers;

/* Expected Output:
| id | name | region |
|----|------|
| 1 | Bob | NA |
| 2 | Alice | EU |
...

*/
-- Basic: No WHERE means all rows
```

• **Ontermediate Example:**

```
-- Retrieve only name and region for users in 'NA'

SELECT name, region

FROM users

WHERE region = 'NA';

/* Expected Output:

| name | region |

|-----|
| Bob | NA |

*/

-- Context: Quick check of user distribution in North America
```

• SRE-Level Example:

```
-- Join users and orders for advanced SRE analysis

SELECT u.name, COUNT(o.order_id) AS total_orders

FROM users u

JOIN orders o ON u.id = o.user_id
```

```
WHERE o.order_date > CURRENT_DATE - INTERVAL '30 days'

GROUP BY u.name

ORDER BY total_orders DESC;

/* Expected Output:

| name | total_orders |

|------|------------|
| Bob | 15 |
| Alice | 10 |

...

*/

-- Production Relevance: Summarize order load per user for trending or anomaly detection
```

Instructional Notes:

- 🖎 **SRE Insight:** Missing a WHERE clause can cause large data scans, spiking I/O usage.
- SRE Insight: Use EXPLAIN or similar tools to check the query plan for performance.
- <u>Common Pitfall:</u> Accidentally writing WHERE region='NA' OR region='EU' incorrectly can lead to unexpected results.
- **Common Pitfall:** Using SELECT * in production.
- **Security Note:** Be mindful not to expose sensitive PII data in logs or shared queries.
- Performance Impact: Proper indexing on columns used in WHERE drastically improves speed.
- Career Risk: A rogue SELECT * from a multi-million-row table can degrade performance and cause an incident.
- Recovery Strategy: If a query is locked or running too long, KILL or CANCEL the session—then investigate root causes.

Command/Concept: Database Connection

Overview:

Establishing a **connection** to the database is your first step before any queries can run.

Real-World Analogy:

Dialing into a **virtual library**—if the door is locked (no credentials), you can't see any books.

Visual Representation:

```
Client/Tool → [Network/Driver] → Database Server

username/password

(or other auth)
```

Syntax & Variations:

Syntax Form	Example	Description	Support/SRE Usage Context
Oracle Connection	sqlplus user/password@hostname:1521/ORCL	Connect using SQL*Plus tool.	Troubleshoot Oracle-based systems
PostgreSQL Connection	psql -h hostname -U user -d dbname	CLI for Postgres.	Basic psql usage.
SQL Server	sqlcmd -S hostname -U user -P password -d dbname	T-SQL CLI.	Windows or cross- platform usage.
MongoDB	<pre>mongo "mongodb://user:pass@host:27017/db"</pre>	Connect to a Mongo instance.	NoSQL environment interactions.
Kafka CLI	kafka-console-consumerbootstrap- server host:9092topic myTopic	Consumes messages from topic.	Streaming logs troubleshooting.

Tiered Examples:

• **@** Beginner Example:

```
# Connect to Postgres
psql -h localhost -U dev_user -d sampledb
# Once connected:
# dev_user@sampledb=>
# You can now run queries like SELECT * FROM customers;
```

• **Ontion** Intermediate Example:

```
# Connect to MongoDB with authentication
mongo "mongodb://alice:secret@prod-mongo:27017/orders_db"
# Then:
# > db.orders.find({ region: "NA" })
# ...
# Cross-Database Variation: Syntax differs from SQL.
```

• SRE-Level Example:

```
# Connect to Kafka console consumer for SRE diagnostics
kafka-console-consumer \
```

```
--bootstrap-server broker1:9092 \
--topic error_logs \
--from-beginning
# Observing streaming logs for real-time incident analysis
```

Instructional Notes:

- **Beginner Tip:** Always verify credentials and connection info. Typos cause repeated connection failures.
- SRE Insight: Monitoring DB connections can reveal environment misconfigurations or capacity issues.
- SRE Insight: Use secure channels (SSL/TLS) in production for privacy and integrity.
- <u>A</u> **Common Pitfall:** Hardcoding credentials in scripts or sharing them openly can lead to security breaches.
- **Common Pitfall:** Using root/admin accounts for everyday queries is risky.
- **Security Note:** Proper role-based access control ensures minimal privileges.
- Performance Impact: Overly frequent connections (lack of pooling) can degrade DB performance.
- Gareer Risk: Using default admin credentials in production is a major vulnerability.
- Recovery Strategy: If locked out, contact DB admins or check backup credentials. Always have an emergency failsafe.

System Effects Section

When you run a **SELECT** query or join multiple tables:

- 1. **Parsing & Optimization**: The database engine checks syntax, picks an execution plan, and possibly uses indexes.
- 2. **Locking/Concurrency**: Read queries might not block others heavily, but large scans can hog CPU or I/O.
- 3. Resource Usage:
 - CPU for parsing and combining data
 - Memory for result caching
 - o I/O to read from disk
 - Network to send results to client
- 4. Monitoring & Observability:
 - Tools like EXPLAIN (Postgres, Oracle, SQL Server) or explain() in Mongo can reveal cost and plan.
 - Kafka queries can be tracked by consumer lag metrics.

Cross-Database Behavioral Differences:

- Oracle, PostgreSQL, SQL Server are similar in how they parse queries but differ in indexing strategies and concurrency models.
- MongoDB can lock at document or collection levels depending on the version and storage engine.
- Kafka is streaming-based; queries equate to reading or filtering event logs.

Warning Signs:

- Long Lock Waits: If your query is blocked or blocking others.
- **High CPU Utilization**: Possibly scanning entire huge tables with no index usage.
- Slow Network Throughput: Returning massive result sets to the client.

SRE Recommendations:

- Use **metrics** like query latency, concurrency levels, and CPU usage.
- Implement alerts if certain queries exceed time thresholds.
- Plan for **scalability** with partitioning or archiving large tables.



😰 Day 1 Visual Learning Aids

- 1. Relational Database Structure: Diagram of schemas, tables, columns, and rows (see example under "Relational Database Fundamentals").
- 2. Primary/Foreign Key Relationship: Visual of two tables connected by a foreign key (shown in the PK/FK section).
- 3. **SQL Query Flow**:

```
- User/Tool -
(e.g., psql)
        \downarrow
   [DB Engine]
   Parsing → Optimization → Execution
        1
    Returns Rows
```

- 4. Cross-Database Syntax Comparison (Table under "SELECT, FROM, WHERE").
- 5. Database Paradigm Comparison:

```
Relational (Tables, Rows, Columns)
NoSQL (Collections, Documents)
Streaming (Topics, Messages)
```

o Explaining how each deals with "data retrieval."



Day 1 Hands-On Exercises

Beginner Exercises (3)

1. Database Connection Exercise

Goal: Connect to a sample database in each system.

Instructions:

- Oracle: sqlplus user/pass@localhost:1521/ORCL
- Postgres: psql -h localhost -U user -d sampledb
- SQL Server: sqlcmd -S localhost -U user -P password -d sampledb
- MongoDB: mongo "mongodb://user:pass@localhost:27017/sampledb"
- Kafka: kafka-console-consumer --bootstrap-server localhost:9092 --topic sampleTopic

Expected Outcome: You can run a basic query or consumer command without errors. **Verification**: You see the DB prompt or streaming data.

2. Basic SELECT Exercise

Goal: Retrieve all rows from a users table/collection/topic.

Instructions:

- Oracle/Postgres/SQL Server: SELECT * FROM users;
- o MongoDB: db.users.find({});
- Kafka: kafka-console-consumer --bootstrap-server ... --topic usersTopic (if messages present)

Expected Outcome: A list of user records is displayed.

Verification: Confirm no errors and correct data format.

3. Simple WHERE Filter Exercise

Goal: Filter users by region (e.g., 'NA').

Instructions:

Oracle/Postgres/SQL Server:

```
SELECT *
FROM users
WHERE region = 'NA';
```

o MongoDB:

```
db.users.find({ region: "NA" })
```

o Kafka (KSQL):

```
SELECT * FROM users_stream
WHERE region = 'NA' EMIT CHANGES;
```

Expected Outcome: Only user rows/docs/messages with region='NA'.

Verification: Check that all returned records match the filter.

Intermediate Exercises (3)

1. Multi-Table Exploration

Goal: Identify relationships between users and orders via primary/foreign keys.

Instructions:

- Inspect table schemas or run DESCRIBE orders; (MySQL/Oracle), \d orders (Postgres).
- Confirm user_id in orders references users(id).

Expected Outcome: Clear understanding of PK/FK relationships.

Verification: You can verbally explain the link between these tables.

2. Column Selection and Filtering

Goal: Write an optimized query returning only name and region from users where region='EU'. **Instructions**:

Example (Postgres/Oracle/SQL Server):

```
SELECT name, region
FROM users
WHERE region = 'EU';
```

o MongoDB:

```
db.users.find({ region: "EU" }, { name: 1, region: 1, _id: 0 });
```

Expected Outcome: Only name and region fields are shown for EU users.

Verification: No extra columns or docs.

3. Support Scenario Query

Goal: Find all orders > \$100 in the last 7 days for a user named "Alice".

Instructions (generic SQL example):

```
SELECT o.order_id, o.order_total
FROM orders o
JOIN users u ON o.user_id = u.id
WHERE u.name = 'Alice'
AND o.order_total > 100
AND o.order_date >= CURRENT_DATE - INTERVAL '7 days';
```

Expected Outcome: Orders matching the filter.

Verification: Ensure the results are accurate and the query runs efficiently.



SRE-Level Exercises (3)

1. Query Performance Analysis

Goal: Examine the execution plan for a SELECT query.

Instructions (Postgres example):

```
EXPLAIN ANALYZE
SELECT *
FROM orders
WHERE order_total > 100;
```

Expected Outcome: Plan details about table scan vs. index usage.

Verification: Identify whether a sequential scan or index scan was used.

2. Cross-Database Query Translation

Goal: Write equivalent queries in Oracle, Postgres, SQL Server, MongoDB, and (optionally) KSQL for retrieving top 5 largest orders.

Instructions:

- Postgres: SELECT * FROM orders ORDER BY order_total DESC LIMIT 5;
- SQL Server: SELECT TOP 5 * FROM orders ORDER BY order_total DESC;
- Oracle (pre-12c): SELECT * FROM (SELECT * FROM orders ORDER BY order total DESC) WHERE ROWNUM <= 5;
- o MongoDB: db.orders.find().sort({order_total: -1}).limit(5);
- KSQL: SELECT * FROM orders_stream PARTITION BY order_total EMIT CHANGES; (top concept might differ)

Expected Outcome: Identical logical result across systems.

Verification: Each approach yields the "top 5" by total.

3. Monitoring Setup

Goal: Configure basic query monitoring for performance tracking.

Instructions:

- Postgres: Enable pg_stat_statements, check pg_stat_activity.
- SQL Server: Use SQL Profiler or Extended Events.
- Oracle: Use AWR or statspack.
- MongoDB: Profiler with db.setProfilingLevel(1).
- Kafka: Monitor consumer lag with kafka-consumer-groups.sh.

Expected Outcome: Observability into who is running slow queries and how often.

Verification: Ensure logs or metrics show up in your chosen dashboards.



Knowledge Check Quiz

Beginner Tier (4 questions)

1. Which statement best describes a 'table' in a relational database?

- A) A random collection of documents
- o B) A structured set of rows and columns
- C) A streaming log partition
- o D) A hidden system object that stores metadata

Explanation: Tables organize data into columns (fields) and rows (records).

2. What does a Primary Key guarantee?

- o A) Multiple rows can share the same ID
- B) Each row is uniquely identified
- o C) Faster writes only
- D) Automatic data archiving

Answer: B

Explanation: PK ensures uniqueness.

3. In SQL, what does SELECT * FROM users; return?

- A) Only rows with null values
- o B) All columns and rows from users
- o C) No rows unless there's a JOIN
- o D) The top 10 rows by default

Answer: B

Explanation: SELECT * returns every column for every row.

4. How would you filter for rows where region='NA'?

```
• A) WHERE region != 'NA'
```

- B) SELECT region = 'NA' FROM table;
- O C) WHERE region = 'NA'
- D) FIND region='NA'

Answer: C

Explanation: The correct condition is WHERE region='NA'.

Intermediate Tier (4 questions)

1. Why is using **SELECT** * in production often discouraged?

- o A) It's shorter to write
- o B) It can return unnecessary columns, causing performance overhead
- o C) Databases prefer returning partial rows for indexing

o D) It's not valid syntax in any database

Answer: B

Explanation: Retrieving unneeded data can slow queries and waste resources.

2. Which is a typical cause of slow queries in a relational database?

• A) Using an appropriate WHERE clause

- B) Proper indexing on frequently filtered columns
- o C) Large table scans due to missing indexes
- o D) Having many small tables

Answer: C

Explanation: Missing indexes lead to full scans, slowing performance.

3. In a typical foreign key relationship, which is true?

- o A) Foreign keys automatically create unique indexes
- o B) Foreign keys link to a primary key in another table
- o C) You only need foreign keys for NoSQL DBs
- o D) A foreign key is the same as a composite key

Answer: B

Explanation: FKs reference a primary key in another table.

4. How might MongoDB handle "joins"?

- A) With the standard SQL JOIN command
- B) Through \$lookup in aggregation or embedded documents
- o C) By referencing a foreign key constraint
- o D) By refusing to combine data from multiple collections

Answer: B

Explanation: \$100kup can simulate joins in MongoDB, but it's not as direct as relational joins.

SRE-Level Tier (4 questions)

1. If a basic SELECT query suddenly takes 10x longer, which factor is most likely at play?

- A) The UI color scheme changed
- o B) An index was dropped or is no longer used
- o C) The table magically lost data
- o D) The data center physically moved

Answer: B

Explanation: Loss of an index or misused index is a common cause of sudden performance hits.

2. Which monitoring metric best indicates heavy read queries in a relational DB?

- o A) CPU load on an unrelated server
- o B) Active connections, read IOPS, or buffer hit ratios
- o C) Network usage on a developer's laptop
- o D) Memory usage on a client machine

Answer: B

Explanation: High read IOPS and buffer usage often signal heavy read queries.

3. In Kafka, retrieving "rows" from a 'topic' is conceptually similar to what in SQL?

- o A) A standard SELECT with a WHERE clause
- o B) Dropping a table
- o C) Creating an index
- o D) Merging partial updates

Answer: A

Explanation: Consuming messages from a topic can be akin to reading rows from a table.

4. How can you mitigate concurrency problems in large relational tables?

- o A) Only let one person query at a time
- B) Use indexing, partitioning, and efficient WHERE clauses
- o C) Randomly kill queries after 10 seconds
- o D) Avoid constraints entirely

Answer: B

Explanation: Good indexing, partition strategies, and careful query design mitigate locking and concurrency issues.

Day 1 Troubleshooting Scenarios

1. Scenario: "Missing Data" Misconception

- **Symptom**: A support analyst claims a user's data is "missing" when they query the database.
- Q Possible Causes:
 - 1. The WHERE clause is too restrictive or incorrect.
 - 2. The user's record exists in a different table, not joined properly.
 - 3. Typo in column/table name leading to no matching results.

• <u>\$\Delta\$</u> Diagnostic Approach:

- 1. Verify user existence with a broad query (SELECT * FROM users;).
- 2. Check table relationships (FK references).
- 3. Compare user ID in both users and orders.

• 🔧 Resolution Steps:

- 1. Correct the WHERE clause to match actual data.
- 2. Use appropriate JOIN if the data is in multiple tables.
- 3. Fix any typos or mismatched columns.
- Prevention Strategy: Understand the schema thoroughly, run smaller test queries first.
- **Knowledge Connection**: Demonstrates the importance of correct SELECT and PK/FK usage.
- **Cross-Database Variation**: Mongo might have data in a different collection. Kafka might require reading from multiple topics.
- **SRE Metrics**: Track query success/failure rates and user error logs for high "no results" queries.

2. Scenario: Slow Query Performance

Symptom: A simple SELECT query is taking 30+ seconds.

Q Possible Causes:

- 1. No WHERE clause, returning massive data sets.
- 2. SELECT * with large, wide tables.
- 3. Missing or dropped index on a frequently searched column.

• <u>\$\Delta\$</u> Diagnostic Approach:

- 1. Check EXPLAIN plan for full table scans.
- 2. Compare performance with a narrower SELECT or indexing.
- 3. Assess concurrency—are multiple heavy queries running?

• 🗞 Resolution Steps:

- 1. Add or refine a WHERE clause.
- 2. Avoid SELECT *; choose specific columns.
- 3. Reintroduce or create relevant indexes.
- **Prevention Strategy**: Regularly tune queries, watch for large data expansions, implement indexing best practices.
- **Knowledge Connection**: Illustrates the cost of poorly structured queries.
- **Cross-Database Variation**: Tools differ—MongoDB uses explain(), Kafka might rely on partition usage metrics.
- **SRE Metrics**: Monitor query response times, concurrency, and CPU usage.

3. Scenario: Connection Issues

- **Symptom**: A new support engineer can't connect to the sample database.
- • Possible Causes:
 - 1. Incorrect credentials or connection string.
 - 2. Firewall or network block.
 - 3. The DB service is offline/not started.

<u>A</u> Diagnostic Approach:

- 1. Verify credentials and host/port.
- 2. Test connectivity with ping or netcat.
- 3. Check DB logs for listener or service errors.

Nesolution Steps:

- 1. Correct the username/password or connection string.
- 2. Open firewall ports or whitelist the IP.
- 3. Start or restart the DB service if down.
- **Prevention Strategy**: Document connection details, ensure minimal but functional privileges.
- **Knowledge Connection**: Ties back to "Database Connection" concept.
- **Cross-Database Variation**: Different ports (1521 for Oracle, 5432 for Postgres, 27017 for Mongo, 9092 for Kafka).
- **IXI** SRE Metrics: Connection success/failure rates, service uptime checks, alerting on DB down states.

? Frequently Asked Questions



- 1. Q: Is SQL case-sensitive?
 - **A**: Many SQL keywords are case-insensitive, but table/column names can be case-sensitive in certain databases (Oracle with quoted identifiers, for example). Always be consistent.
- 2. **Q**: Can I learn databases if I'm not a developer?
 - **A**: Absolutely! Understanding data retrieval is crucial for support tasks. You'll quickly see how queries speed up troubleshooting.
- 3. Q: Why do I need a primary key in every table?
 - A: It ensures each record is uniquely identifiable, preventing duplicates and helping with fast lookups.

Intermediate FAQs

- 1. **Q**: When should I use a composite primary key?
 - **A**: If you need to uniquely identify a row by more than one column (e.g., user_id + group_id in a membership table). This is common in many-to-many relationships.
- 2. **Q**: How do I speed up queries that filter on the same column frequently?
 - **A**: Creating an index on that column typically speeds up lookups significantly, especially if the table is large.
- 3. **Q**: What's the difference between Oracle's ROWNUM and SQL Server's TOP or Postgres' LIMIT? **A**: They are different ways to limit query results. Oracle uses ROWNUM (or FETCH FIRST n ROWS ONLY in newer versions), SQL Server uses TOP, and Postgres uses LIMIT.

SRE-Level FAQs

- 1. **Q**: How do I handle large tables that keep growing?
 - **A**: Strategies include partitioning, archiving old data, and implementing efficient indexing. Monitoring size growth is crucial for capacity planning.
- 2. **Q**: What are typical concurrency issues I might see?
 - **A**: Row locking, deadlocks, or high CPU usage from multiple heavy queries. Tools like "deadlock graphs" in SQL Server or Postgres logs help debug.
- 3. **Q**: How does Kafka fit with relational databases?
 - **A**: Kafka can stream real-time data that you might later store in a relational database for analysis. Or you can replicate DB changes into Kafka for event-driven architectures.



) Support/SRE Scenario

Scenario: Production login service is slow, and you suspect the user lookup queries are part of the problem.

Steps (6 total):

1. Identify Query

```
SELECT *
FROM users
WHERE username = 'alice';
```

• **Reasoning**: See if this query runs often and how it performs.

2. Check Execution Plan

```
EXPLAIN ANALYZE

SELECT *

FROM users

WHERE username = 'alice';
```

• Reasoning: Determine whether it's using an index or scanning the entire table.

3. Evaluate Index

• Reasoning: If username isn't indexed, the DB might do a full scan for each login attempt.

4. Create or Fix Index

```
CREATE INDEX idx_users_username
ON users(username);
```

Reasoning: Speed up lookups by username. For large user tables, this is crucial.

5. Monitor Query Speed

• **Reasoning**: Confirm performance improvements using DB logs or external metrics.

6. Kafka Integration (optional)

 If login attempts also produce Kafka messages, ensure consumer lag isn't impacting real-time analysis.

Connection to SRE Principles:

- Observability: Use logs/metrics to track query times.
- Reliability: Indexing ensures stable performance under load.
- Minimal Downtime: Proactive measures to avoid future slowdowns.

Key Takeaways

1. 5+ Command/Concept Summary Points:

1. Tables, Columns, Rows form the backbone of relational data structures.

- 2. **Primary Keys** uniquely identify each record, **Foreign Keys** link tables together.
- 3. **SELECT-FROM-WHERE** is the fundamental triad for retrieving data.
- 4. **Database Connections** vary by platform (Oracle, Postgres, SQL Server, MongoDB, Kafka), but the idea remains the same.
- 5. Cross-Database Variation means syntax and feature differences, but principles stay consistent.

2. 3+ Operational Insights (Reliability):

- Using correct PK/FK constraints ensures data consistency.
- Monitoring queries prevents unplanned downtime from runaway operations.
- Index usage drastically affects performance and concurrency under load.

3. 3+ Best Practices (Performance):

- Avoid SELECT * in production; specify columns.
- Always test queries on dev/staging to confirm performance and correctness.
- Keep an eye on table/column usage and add or adjust indexes as data grows.

4. 3+ Critical Warnings/Pitfalls:

- Missing WHERE clause can blow up resource usage.
- Not having a PK leads to data duplication.
- Hardcoding credentials or ignoring security best practices can lead to breaches.

5. 3+ Cross-Database Considerations:

- Oracle might need more formal quoting for identifiers, while SQL Server uses [].
- MongoDB doesn't have strict PK/FK but uses <u>id</u> and references.
- Kafka "queries" revolve around streaming data consumption, not structured tables.

6. 3+ Monitoring Recommendations:

- Use EXPLAIN or ANALYZE to monitor query plans.
- Track concurrency, CPU usage, memory, and I/O for suspicious spikes.
- Implement basic alerting on query latency and error rates.

Connections to Support/SRE Excellence: Understanding these fundamental concepts ensures you respond to incidents swiftly, design queries with reliability in mind, and avoid data disasters from day one.



Day 1 Career Protection Guide

High-Risk SELECT Operations

1. **SELECT * on Very Large Tables**

- Real Incident: A single SELECT * FROM logs; crashed production due to huge I/O spikes.
- Warning Signs: Overly high CPU/disk usage.

2. Missing WHERE Clause

- Real Incident: Query returned millions of rows to the client, saturating network.
- Warning Signs: Long run times, "hangs," or out-of-memory errors.

3. Joining Many Tables with No Proper Index

- Real Incident: Complex JOIN across multiple big tables timed out, blocking other queries.
- Warning Signs: Surging concurrency, lock waits.

Verification Best Practices

- LIMIT/TOP usage to test queries before retrieving entire dataset.
- **Test in Dev** with sample data first, especially if the query can be large.
- Check Execution Plans for suspicious full scans or missing indexes.

Recovery Strategies

- Cancel or Kill the session if a query is locked or hogging resources.
- Communicate promptly if you suspect your query caused an incident.
- Roll Back partial changes if your query had write operations (in future sessions).

Cross-Database Considerations

- Oracle: Resource Manager might throttle extreme queries.
- PostgreSQL: pg_cancel_backend() can kill a bad query.
- **SQL Server**: Use Management Studio to kill spid.
- MongoDB: kill0p() to terminate a long-running operation.
- **Kafka**: Large queries with KSQL might stress the broker if partitions are huge.

First-Day Safeguards

- Access Control: Start with read-only roles.
- Query Reviews: Ask a senior dev/DBA to review suspicious or complex queries.
- Small Batches: Retrieve data in smaller chunks.
- Backups: Confirm a restore plan is in place before running any critical queries.



Preview of Next Topic

On Day 2, we'll focus on CRUD Operations (INSERT, UPDATE, DELETE) and the fundamentals of transaction control. You'll see how to alter data safely—especially important for reliability and preventing data corruption. We'll also introduce how NoSQL databases handle writes differently and how Kafka stream processing can feed or consume data updates in real time.

Preparatory Suggestions:

- Practice writing selective queries with small result sets.
- Experiment with indexing a test table to see performance differences.
- Familiarize yourself with your environment's rollback/commit process.

Day 1 Further Learning Resources

Beginner SQL & Relational Database Resources (3)

1. SQLBolt

- Link: https://sqlbolt.com/
- **Description**: Interactive lessons on SELECT, FROM, WHERE, perfect for newcomers.
- Helps Support Roles: Short tutorials that build confidence in basic SQL queries.
- **Estimated Time**: 1-2 hours to complete initial lessons.

2. Khan Academy: Intro to SQL

- **Link**: https://www.khanacademy.org/computing/computer-programming/sql
- **Description**: Friendly approach with examples and exercises on fundamental queries.
- Helps Support Roles: Visual demos for queries, suitable for those with no DB experience.
- Estimated Time: 2-3 hours for the SQL unit.

3. W3Schools SQL Tutorial

- Link: https://www.w3schools.com/sql/
- Description: Provides an overview of SQL syntax with interactive try-it-yourself.
- Helps Support Roles: Quick references for basic commands.
- Estimated Time: 1 hour to browse main topics.

Intermediate Relational Concepts Resources (3)

1. PostgreSQL Official Documentation

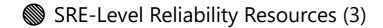
- Link: https://www.postgresql.org/docs/
- **Description**: Deep dive on table relationships, constraints, and indexing.
- **Practical Connection**: Builds upon Day 1 by exploring more advanced relational features.
- Key Takeaway: Understand performance implications of schema design.

2. Microsoft SQL Server Tutorials

- Link: https://learn.microsoft.com/en-us/sql/
- **Description**: Covers T-SQL queries, advanced SELECT variations, indexing strategies.
- **Practical Connection**: Helps you apply queries to real support tasks with SQL Server.
- Key Takeaway: Leverages built-in tools like Query Analyzer for performance checks.

3. Oracle Live SOL

- Link: https://livesql.oracle.com/
- **Description**: Free environment to practice SQL with Oracle-specific features.
- **Practical Connection**: Focuses on constraints, indexes, and more advanced queries.
- Key Takeaway: Hands-on labs for developing robust Oracle-based solutions.



1. Database Reliability Engineering (O'Reilly)

- Link: https://www.oreilly.com/library/view/database-reliability-engineering/9781491925935/
- Description: Connects fundamental SQL knowledge to broader reliability topics.
- SRE Impact: Emphasizes backups, HA, performance monitoring.
- Value: Ideal for bridging Day 1 basics to advanced SRE roles.

2. High Performance MySQL

- Link: https://www.oreilly.com/library/view/high-performance-mysql/9781449332471/
- **Description**: While MySQL-specific, the performance principles apply across relational systems.
- **SRE Impact**: Deep dives on indexing, query optimization, and replication.
- Value: Enhances reliability perspective even if you're on Postgres, Oracle, or SQL Server.

3. Kafka: The Definitive Guide

- Link: https://www.confluent.io/resources/kafka-the-definitive-guide/
- **Description**: Explains Kafka's streaming model, how it differs from relational DBs.
- **SRE Impact**: Illustrates reliability patterns for high-velocity data streams.
- **Value**: Understanding event streaming complements relational knowledge for real-time systems.

Database-Specific Resources for Day 1 Concepts (5)

- 1. Oracle: https://docs.oracle.com/en/database/
 - Value: Official Oracle Database documentation for beginners on SELECT queries, connecting to a sample DB.
 - **Tutorial**: Oracle "2 Day DBA" for step-by-step environment setup.
- 2. **PostgreSQL**: https://www.postgresql.org/docs/current/tutorial-start.html
 - **Value**: Basic tutorial on table creation, querying, and psql usage.
 - **Tutorial**: Includes examples for creating a test database and running queries.
- 3. **SQL Server**: https://learn.microsoft.com/en-us/sql/ssms/quickstart-sql-server-management-studio
 - Value: Quickstart guide for using SQL Server Management Studio to connect and query.
 - Tutorial: Step-by-step on installing SSMS and running your first SELECT statement.
- 4. MongoDB: https://docs.mongodb.com/manual/core/databases-and-collections/
 - **Value**: Explains how "collections" and "documents" compare to relational tables/rows.
 - Tutorial: Create a sample database, insert documents, and query them.
- 5. **Kafka**: https://kafka.apache.org/quickstart
 - Value: Simple tutorial for running Kafka locally, creating a topic, and producing/consuming messages.
 - o **Tutorial**: Demonstrates how streaming data operations differ from SQL queries.



Closing Message

Congratulations on completing Day 1 of our SRE Database Training Module! You've gained a firm grasp of tables, columns, rows, how primary and foreign keys maintain integrity, and how to perform basic SELECT queries across multiple databases. By embracing these fundamentals with an SRE mindset, you're already on track to write safer queries, diagnose issues more effectively, and help maintain reliable systems.

Next Steps:

- Experiment with the recommended resources and labs.
- Practice writing targeted SELECT queries in dev environments.
- Prepare for Day 2, where we'll explore CRUD operations, transactions, and advanced reliability strategies.

Keep these basics close at hand, and remember: a well-structured query is your best friend in diagnosing and preventing issues. Enjoy your journey toward SRE excellence!