SRE Database Training Module - Day 1: Core Relational Database Fundamentals (Oracle Focus)

Welcome to Day 1 of our SRE Database Training module! Today we'll establish the foundation of relational database concepts using **Oracle** as our primary example, with key comparisons to **PostgreSQL** and **SQL Server**. We'll build knowledge "brick by brick," from absolute basics to SRE-level insights. By the end of this module, you'll have a solid understanding of core database structures and operations, and be ready to tackle more advanced topics in subsequent days.

Introduction

Relational databases are the backbone of reliable systems in the enterprise world. Whether supporting a mission-critical ERP application or a smaller application's backend, databases store and organize data to be retrieved and manipulated efficiently.

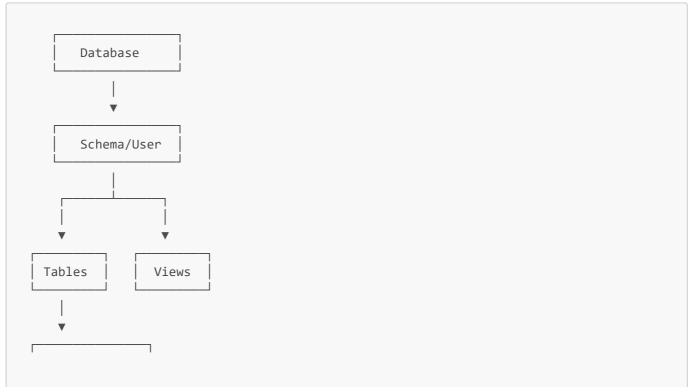
Why Oracle?

Oracle Database remains one of the most widely used and powerful relational database systems in enterprise environments. Its advanced features, comprehensive toolset, and robust performance tuning capabilities make it a common choice where reliability and scale are paramount.

In today's session, you will:

- 1. Explore fundamental database concepts (tables, columns, rows, primary keys, foreign keys).
- 2. **Learn** and practice basic SQL queries (SELECT, FROM, WHERE) using Oracle syntax.
- 3. Compare Oracle syntax to PostgreSQL and SQL Server to broaden your skillset.
- 4. **Understand** how these concepts are used daily in Product Support and SRE roles.

Below is a high-level concept map to visualize the relationships among fundamental components:



Rows/Columns

Oracle is our primary focus, but key differences in PostgreSQL and SQL Server will be shown along the way to help you work with various systems.

& Learning Objectives by Tier

Below are four measurable objectives at each tier—beginner, intermediate, and SRE-level—to guide your progress.

- Beginner Objectives
 - 1. Identify the basic structures of a relational database (tables, rows, columns).
 - 2. Explain the concept of primary and foreign keys.
 - 3. **Construct** simple SELECT queries using FROM and WHERE.
 - 4. Connect to an Oracle database via SQL*Plus or SQL Developer.
- Intermediate Objectives
 - 1. Differentiate Oracle syntax from PostgreSQL and SQL Server for common SQL statements.
 - 2. Analyze basic troubleshooting steps for Oracle (e.g., invalid username/password, locked account).
 - 3. **Apply** indexing and constraints to improve data integrity and performance.
 - 4. **Use** Oracle data dictionary views (e.g., ALL_TABLES, ALL_CONSTRAINTS) to gather metadata.
- SRE-Level Objectives
 - 1. **Demonstrate** performance monitoring using Oracle's dynamic performance views (e.g., V\$SESSION, V\$SQL).
 - 2. Interpret Oracle execution plans for query optimization.
 - 3. **Implement** Oracle-specific recovery strategies (e.g., Flashback, RMAN).
 - 4. **Correlate** database events to SRE principles like reliability, observability, and scalability.

Core Concepts

This module uses a brick by brick approach for each core concept. You'll see:

- Beginner Analogy to give a simple real-world connection.
- **Visual Representation** to illustrate structures.
- <u>A</u> **Technical Explanation** of the concept.
- Support/SRE Application to show how it impacts your daily tasks.
- System Impact on performance, reliability, and scalability.
- **Common Misconceptions** to help you avoid pitfalls.
- IJ SQL Dialect Comparison to highlight Oracle, PostgreSQL, and SQL Server differences.

We'll focus on **Relational Database Structure**, **Primary/Foreign Keys**, **Basic SELECT Statement**, **FROM Clause**, and **WHERE Clause**.

Day 1 Concept Breakdown

1. Relational Database Structure

Concept Overview:

A relational database organizes data into tables, which have rows (records) and columns (attributes). Each column has a data type, and each row represents one instance of whatever entity the table tracks.

Beginner Analogy:

Think of a spreadsheet with columns labeled (Name, Address, etc.) and each row containing a single record (a person's information).

☑ Visual Representation (ASCII Diagram):

EMD TD	ETDET NAME	LACT NAME
EMP_ID	FIRST_NAME	LAST_NAME
1001	John	Doe
1002	Jane	Smith
1003	Alice	Johnson

<u>A</u> Technical Explanation:

- **Table**: A collection of rows sharing the same columns.
- Row (record): A single entry in the table.
- Column (field): An attribute or property of the entity.

Support/SRE Application:

- Each table underpins specific functionality (customers, orders, etc.). Knowing how data is structured helps diagnose issues when queries fail or performance slows.
- Understanding table design is crucial for identifying potential data integrity issues or performance bottlenecks.

System Impact:

- Proper normalization (organizing columns logically to reduce redundancy) enhances performance and maintainability.
- Poor structure leads to data anomalies and slow gueries.

⚠ Common Misconceptions:

Misconception: "Storing everything in one big table is easiest."
 Correction: This can cause massive redundancy and performance issues. Splitting data across multiple related tables is best practice.

III SQL Dialect Comparison:

Operation	Oracle Example	PostgreSQL Example	SQL Server Example	Notes
Create a	CREATE TABLE	CREATE TABLE	CREATE TABLE	Syntax is similar across all
	employees (employees (employees (three, minor data type
table);););	differences.

Operation	Oracle Example	PostgreSQL Example	SQL Server Example	Notes
List tables	SELECT table_name FROM user_tables;	<pre>\dt (psql command) or SELECT tablename;</pre>	SELECT name FROM sys.tables;	Oracle uses data dictionary views; PG/SQL Server use system schemas.
Describe table	DESC employees; (SQL*Plus)	<pre>\d employees (psql) or SELECT column_name;</pre>	In SSMS, expand table or sp_columns employees;	Each system has its own shorthand or stored procedures to show schema.

2. Primary Keys and Foreign Keys

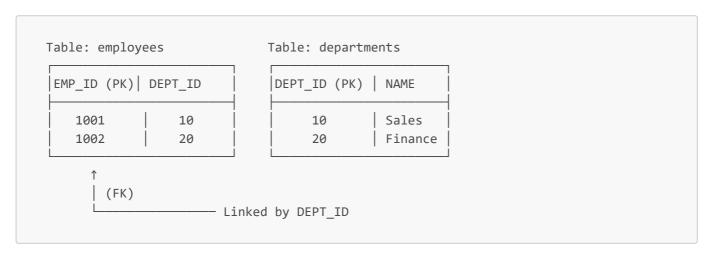
Concept Overview:

A **Primary Key (PK)** uniquely identifies each row in a table. A **Foreign Key (FK)** creates relationships between tables, linking a column in one table to a primary key in another.

Beginner Analogy:

- **Primary Key**: Your "Social Security Number" or "Passport Number" a unique ID that identifies you.
- **Foreign Key**: A reference link. If you have an "Address" table, each address row might reference the unique ID from the "Person" table.

☑ Visual Representation (ASCII Diagram):



A Technical Explanation:

- Primary Key: Implies uniqueness and often not-null constraints.
- Foreign Key: Enforces referential integrity; a dept_id in employees must exist in departments.

Support/SRE Application:

• Quick detection of data issues: if a foreign key constraint fails, your troubleshooting might begin with verifying data consistency between linked tables.

System Impact:

 Foreign keys maintain data integrity but can slow down large INSERT or UPDATE operations if not indexed properly.

⚠ Common Misconceptions:

Misconception: "Foreign keys are optional overhead."
 Correction: They're essential for data consistency, preventing orphan records.

SQL Dialect Comparison:

Operation	Oracle Example	PostgreSQL Example	SQL Server Example	Notes
Create PK	CREATE TABLE departments (dept_id NUMBER PRIMARY KEY,);	CREATE TABLE departments (dept_id INT PRIMARY KEY,);	CREATE TABLE departments (dept_id INT PRIMARY KEY,);	Very similar, data types differ (NUMBER vs INT).
Create FK	CREATE TABLE employees (dept_id NUMBER REFERENCES departments(dept_id));	CREATE TABLE employees (dept_id INT REFERENCES departments(dept_id),);	CREATE TABLE employees (dept_id INT FOREIGN KEY REFERENCES departments(dept_id),);	SQL Server often uses a separate ALTER TABLE ADD CONSTRAINT approach, but inline also works.
Add PK to existing table	ALTER TABLE employees ADD PRIMARY KEY(emp_id);	ALTER TABLE employees ADD PRIMARY KEY(emp_id);	ALTER TABLE employees ADD PRIMARY KEY(emp_id);	Similar syntax across all.
Add FK to existing table	ALTER TABLE employees ADD CONSTRAINT fk_dept FOREIGN KEY(dept_id) REFERENCES departments(dept_id);	ALTER TABLE employees ADD CONSTRAINT fk_dept FOREIGN KEY(dept_id) REFERENCES departments(dept_id);	ALTER TABLE employees ADD CONSTRAINT fk_dept FOREIGN KEY(dept_id) REFERENCES departments(dept_id);	Also quite consistent across systems.

3. Basic SQL SELECT Statement

Concept Overview:

The SELECT statement retrieves data from tables. It typically includes SELECT <columns>, FROM , and optional clauses like WHERE, GROUP BY, and ORDER BY.

Beginner Analogy:

Imagine searching for specific files in a folder. SELECT is like specifying which attributes of the file you need, and FROM is the folder name. WHERE is how you filter results.

☑ Visual Representation (High-Level Query Flow):

```
| SELECT columns |
```

```
FROM table

| WHERE conditions
| GROUP BY ...
| ORDER BY ...

↓

Database Engine

↓

Filtered Results
```

<u>A</u> Technical Explanation:

- **SELECT**: Lists which columns or expressions to return.
- **FROM**: Specifies the table(s) to query.
- Additional clauses refine or group the data.

Support/SRE Application:

- When troubleshooting data issues, you'll often run SELECT queries with filters to confirm if data is correct.
- SELECT is the foundation of all read queries, so performance often matters (e.g., indexing, execution plans).

System Impact:

• Poorly optimized SELECT queries can hog system resources, degrade performance, and cause slowdowns.

⚠ Common Misconceptions:

Misconception: "SELECT * is harmless."
 Correction: It can be expensive, retrieving unneeded columns and causing performance issues.

SQL Dialect Comparison:

Operation	Oracle Example	PostgreSQL Example	SQL Server Example	Notes
Basic SELECT	SELECT emp_id, first_name FROM employees;	<pre>SELECT emp_id, first_name FROM employees;</pre>	<pre>SELECT emp_id, first_name FROM employees;</pre>	Largely identical for basic queries across all three.
Limiting rows	SELECT * FROM employees WHERE ROWNUM <= 5;	SELECT * FROM employees LIMIT 5;	SELECT TOP 5 * FROM employees;	Oracle uses ROWNUM, PostgreSQL uses LIMIT, SQL Server uses TOP.
Using Aliases	SELECT e.emp_id AS ID FROM employees e;	SELECT e.emp_id AS ID FROM employees e;	SELECT e.emp_id AS ID FROM employees e;	Similar, with minor variations in older SQL Server versions.
Concatenation	<pre>SELECT first_name ' ' last_name FROM employees;</pre>	<pre>SELECT first_name ' ' last_name FROM employees;</pre>	<pre>SELECT first_name + ' ' + last_name FROM employees;</pre>	Oracle/PostgreSQL use , SQL Server uses +.

4 FROM Clause and Table Selection

Concept Overview:

FROM indicates which table(s) the database engine should read. Queries can join multiple tables using relational links.

Beginner Analogy:

If you're looking for specific books, the FROM clause is like specifying which bookshelf (table) you're searching.

☑ Visual Representation (Join Example):

```
SELECT e.emp_id, e.first_name, d.name
FROM employees e
JOIN departments d ON e.dept_id = d.dept_id;
```

```
employees (e) departments (d)

[emp_id] name

↑ ↑

joined by dept_id joined by dept_id
```

<u>A</u> Technical Explanation:

- Single Table: FROM <table_name>
- Multiple Tables (JOIN): Different join types (INNER, LEFT, RIGHT, FULL) specify how data aligns.

Support/SRE Application:

• Incorrect FROM or join conditions lead to incomplete or incorrect data, a common cause of application bugs.

System Impact:

• Joins can become expensive with large tables; indexing on join columns is vital for performance.

⚠ Common Misconceptions:

Misconception: "Outer joins always return more rows."
 Correction: Outer joins fill in missing data for the "outer" side, but the final row count depends on the data distribution.

III SQL Dialect Comparison:

Operation	Oracle Example	PostgreSQL Example	SQL Server Example	Notes
Select single table	SELECT * FROM employees;	SELECT * FROM employees;	SELECT * FROM employees;	Identical.

Operation	Oracle Example	PostgreSQL Example	SQL Server Example	Notes
Inner join	SELECT e.emp_id, d.name FROM employees e JOIN departments d ON e.dept_id = d.dept_id;	SELECT e.emp_id, d.name FROM employees e JOIN departments d ON e.dept_id = d.dept_id;	SELECT e.emp_id, d.name FROM employees e JOIN departments d ON e.dept_id = d.dept_id;	Standard ANSI join syntax is the same across all.
Left outer join	SELECT e.emp_id, d.name FROM employees e LEFT JOIN departments d ON;	SELECT e.emp_id, d.name FROM employees e LEFT JOIN departments d ON;	SELECT e.emp_id, d.name FROM employees e LEFT JOIN departments d ON;	Works similarly across Oracle, PG, and SQL Server.
Old Oracle join syntax	SELECT e.emp_id, d.name FROM employees e, departments d WHERE e.dept_id = d.dept_id;	n/a (not recommended in modern usage)	n/a	Oracle had older (+) syntax for outer joins, but it's mostly replaced by standard ANSI joins.

5. WHERE Clause and Basic Filtering

Concept Overview:

The WHERE clause filters rows based on conditions. Only rows meeting those conditions are returned.

Beginner Analogy:

When searching for files in a folder by name or extension, you use a filter (WHERE) to only see matching files.

☑ Visual Representation (Filter Flow):

A Technical Explanation:

- WHERE filters rows before grouping or ordering.
- Logical operators include =, <, >, >=, LIKE, IN, etc.

Support/SRE Application:

- Pinpoint the specific rows that cause an application error or slow performance.
- Filter out extraneous data for targeted analysis.

System Impact:

• Proper use of indexes on the filtered columns can drastically improve query performance.

⚠ Common Misconceptions:

Misconception: "WHERE 1=1 is needed for queries."
 Correction: This is sometimes used in dynamic SQL generation, but not mandatory for normal queries.

SQL Dialect Comparison:

Operation	Oracle Example	PostgreSQL Example	SQL Server Example	Notes
Basic filter	SELECT * FROM employees WHERE dept_id = 10;	SELECT * FROM employees WHERE dept_id = 10;	SELECT * FROM employees WHERE dept_id = 10;	Identical for basic equality filters.
Wildcard (string match)	SELECT * FROM employees WHERE last_name LIKE 'S%';	<pre>SELECT * FROM employees WHERE last_name LIKE 'S%';</pre>	SELECT * FROM employees WHERE last_name LIKE 'S%';	Also very similar across all systems.
Range filter	SELECT * FROM employees WHERE emp_id BETWEEN 100 AND 110;	SELECT * FROM employees WHERE emp_id BETWEEN 100 AND 110;	SELECT * FROM employees WHERE emp_id BETWEEN 100 AND 110;	Uniform usage of BETWEEN.
Set membership	SELECT * FROM employees WHERE dept_id IN (10, 20);	SELECT * FROM employees WHERE dept_id IN (10, 20);	SELECT * FROM employees WHERE dept_id IN (10, 20);	Same usage for IN.

% Oracle-Specific Tools and Techniques

Oracle offers several tools for interacting with and examining database objects:

- 1. **SQL*Plus** A command-line interface for executing SQL and PL/SQL.
- 2. **SQL Developer** A graphical interface for running queries, browsing schema objects, and debugging.
- 3. Enterprise Manager (OEM) A web-based console for monitoring and administering Oracle databases.
- 4. **Oracle Data Dictionary Views** E.g. USER_TABLES, ALL_TABLES, USER_OBJECTS, V\$SESSION, V\$SQL to view metadata and performance metrics.

Use these tools to:

- Check table and index structures (DESC, USER_INDEXES).
- Investigate performance or connectivity issues using OEM or dictionary views.
- Monitor sessions and resource usage (V\$SESSION, V\$SQLAREA).

Oracle Performance Monitoring and Execution Plans

Understanding how Oracle executes a query is crucial for SRE roles:

• Generating Execution Plans

- Use EXPLAIN PLAN FOR <query>; SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
- Or AUTOTRACE in SQL*Plus.

• Key Oracle Performance Views

- V\$SESSION shows active sessions.
- V\$SQL shows recent SQL statements.
- V\$SQLAREA aggregates statements to check overall performance.

• Performance Monitoring

- Tools like Enterprise Manager or TOAD can visualize execution plans.
- o Index usage, table scans, and join methods are highlighted in the plan.

Hands-On Exercises

Beginner Exercises

1. Create and Query a Table

• Create a simple departments table in Oracle. Insert 3 rows. Write a SELECT to retrieve all rows.

2. Describe Table

• Use DESC in SQL*Plus or use Oracle SQL Developer to view columns.

3. Basic Filter

Write a query to return only rows where department_name is 'Sales'.

Intermediate Exercises

1. Primary/Foreign Key Relationship

Create an employees table referencing the departments table. Confirm foreign key constraints.

2. Joining Tables

Write a query joining employees and departments. Filter by a specific dept_id.

3. Comparing Dialects

 Write an equivalent SELECT in Oracle, PostgreSQL, and SQL Server that retrieves employee IDs, first names, and department names.

SRE-Level Exercises

1. Execution Plan Analysis

o Generate an execution plan for a join query and identify whether an index is used.

2. Performance Check

 Inspect the V\$SQLAREA or V\$SQL views to see how many times your query has been executed and the average elapsed time.

3. Optimize Query

 Add an index to improve performance of a slow query. Compare the execution plan before and after indexing.

Troubleshooting Scenarios

1. Locked Account

- **Symptoms**: User cannot log in; receives ORA-28000 error.
- **Cause**: Database user locked after too many failed attempts.
- **Diagnostic**: Check DBA_USERS or attempt to CONNECT.

Resolution: ALTER USER <username> ACCOUNT UNLOCK;

2. Missing Rows in Join

- Symptoms: Query results are incomplete when joining multiple tables.
- Cause: Department ID in employees table does not match any ID in departments.
- Diagnostic: Check foreign key integrity or use an outer join to see missing associations.
- Resolution: Correct data or insert missing department row.

3. Slow Query

- **Symptoms**: A simple **SELECT** takes too long.
- Cause: Full table scan on a large table, no index on the filtering column.
- **Diagnostic**: Use EXPLAIN PLAN or Enterprise Manager to confirm table scan.
- **Resolution**: Create an index, gather statistics, or rewrite the query.

? Frequently Asked Questions

Beginner FAQs

1. How do I log in to Oracle with SQL*Plus?

• sqlplus username/password@hostname:port/SID

2. What's the difference between a database and a table?

 A database is a collection of related tables, indexes, and other objects. A table is just one object storing rows/columns.

3. How do I list my tables?

SELECT table_name FROM user_tables;

Intermediate FAQs

1. How do I see constraints on a table?

```
    SELECT constraint_name, constraint_type FROM user_constraints WHERE
table_name='EMPLOYEES';
```

2. Why does my foreign key insertion fail?

The referenced primary key might not exist in the parent table. Check data integrity.

3. Is it safe to use SELECT *?

Generally no. It can impact performance and return unwanted columns.

SRE-Level FAOs

1. How do I generate an execution plan quickly?

• EXPLAIN PLAN FOR <query>; SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);

2. Which views show long-running sessions?

• V\$SESSION and V\$SESSION_LONGOPS can help identify sessions running for a long time.

3. What's the best way to monitor concurrency in Oracle?

• Use Oracle's Automatic Workload Repository (AWR) reports or V\$ACTIVE_SESSION_HISTORY to track wait events and concurrency issues.

Oracle-Specific SRE Scenario

Incident: Slow Response Times During Peak Hours

Context: A production environment experiences performance degradation at 9 AM every day.

1. **Symptom**: Users complain that logging in and retrieving records is very slow (~30 seconds).

2. Initial Check:

- Check V\$SESSION for active sessions.
- Discover a large batch job is running concurrently.

3. Diagnosis:

- o Generated an execution plan for the main user query. Found a full table scan on a 5-million-row table.
- Found that the batch job locks some rows and causes heightened I/O.

4. SRE Approach:

- Confirm row-level locking is being held by the batch job using V\$LOCK or DBA_BLOCKERS.
- o Check indexing strategy for the user query's WHERE clause.
- Inspect AWR for historical performance trends.

5. Resolution:

- Add an index on the commonly filtered column to avoid full scans.
- Reschedule the batch job to a less congested time.
- o Confirm improvement by re-checking execution plan and system metrics.

6. Outcome:

- User query time reduced to <1 second.
- No further collisions with batch job.

This scenario illustrates how Oracle-specific views (V\$SESSION, V\$LOCK, AWR) and tools (execution plans) help SREs diagnose and fix performance bottlenecks.

Key Takeaways

- Relational Foundations: Tables, columns, rows, keys, and how they interrelate.
- **Oracle-Specific Focus**: Tools like SQL*Plus, SQL Developer, data dictionary views, and Oracle's dynamic performance views provide critical insights.
- Performance: Indexing and understanding execution plans are essential SRE skills.
- **Reliability**: Proper constraints (PK, FK) protect data integrity and reduce the risk of silent data errors.
- **Troubleshooting**: Systematic diagnosis using dictionary views and joins clarifies data mismatches and performance issues.
- **Career Protection**: Stay cautious with high-risk operations (locking, recovery, dropping objects), and always have a rollback or backup plan.

▲ Oracle Career Protection Guide

1. High-Risk Operations

Dropping or truncating tables. Always confirm with multiple levels of backup or approvals.

2. Recovery Strategies

- Use **Flashback** to revert to a previous state if you accidentally remove/modify data.
- RMAN backups ensure you can recover from catastrophic data loss or corruption.

3. Verification Best Practices

- Always test queries in a non-production environment.
- Use read-only queries or placeholders to confirm expected results before running them in production.

4. First-Day Safeguards

- o Review user privileges.
- o Confirm environment (Dev vs. Production) to avoid accidental changes.

Preview of Next Day's Content

On Day 2, we'll dive deeper into:

- Advanced JOIN Techniques (subqueries, common table expressions).
- Data Manipulation (INSERT, UPDATE, DELETE) with Oracle best practices.
- Transaction Management (ACID properties, COMMIT, ROLLBACK).
- Intro to Oracle Index Types and partitioning for large tables.
- More Performance Tuning with real-world SRE examples.

Stay curious and get ready to continue your journey into the world of Oracle database operations!

Congratulations on completing Day 1 of the SRE Database Training module! You have built a solid foundation of **relational database fundamentals**, learned **basic Oracle syntax**, and explored key **SRE-oriented troubleshooting** methods. Continue practicing, and see you on Day 2 for more advanced topics.